

CSE 713: Artificial Intelligence

Search Algorithms — Complete Model Answers

University of Chittagong — 7th Semester B.Sc. (Engg.) Examination

Covers Years: 2024, 2023, 2022, 2021, 2020 — Topics: UCS, BFS, DFS, Greedy, A*,
IDDFS, IDA*

Contents

1	Heuristic Search — Theoretical Foundations	3
1.1	BFS (Uninformed) vs. Best-First Search (Informed)	3
1.2	Informed vs. Uninformed Search — The Fundamental Divide	5
1.3	Progression of Uninformed Search Algorithms	6
1.4	Heuristic Search — Introduction & the Best-First Family	10
1.5	The Heuristic Function $h(n)$	11
1.6	Greedy Best-First Search — Algorithm, Example, Properties	12
1.7	Greedy vs. A* — Direct Comparison	14
2	Examination 2024 — Q2 (9 Marks)	16
2.1	Problem Data	16
2.2	Part a(i) — State-Space Graph [1 mark]	17
2.3	Part a(ii) — Search Algorithm Traces [5 marks]	18
2.4	Part b — Problem Formulations [3 marks]	23
3	Examination 2023 — Q2 (9 Marks)	25
3.1	Problem Data	25
3.2	Part a — Problem Formulations [2 marks]	25
3.3	Part b(i) — State-Space Diagram [0.5 mark]	26
3.4	Part b(ii) — Search Algorithm Traces [4.5 marks]	27
3.5	Part c — IDDFS: Synthesising Benefits of BFS and DFS [2 marks]	30
4	Examination 2022	33
4.1	Section A — Q.A-4(a): DFS vs BFS [2 marks]	33
4.2	Section B — Q.B-4(b): Graph Trace [3 marks]	36
5	Examination 2021	37
5.1	Q2(a) — State Space Search Problem [0.75 mark]	37
5.2	Q2(b) — Problem Formulations [3 marks]	37
5.3	Q2(c) — Graph Traces [5 marks]	39
5.4	Q3(b) — DFID vs IDA* [1.75 marks]	42

6	Examination 2020	45
6.1	Q2(a) — Formulating a Search Problem [1.75 marks]	45
6.2	Q2(b) — Problem Formulations [2 marks]	46
6.3	Q2(c) — Graph Traces [5 marks]	47
6.4	Q3(b) — IDDFS Combines BFS + DFS [2 marks]	48
6.5	Q3(c) — BFS Admissibility [1.75 marks]	50
6.6	Q7(b) — A* Heuristic: Underestimate vs Overestimate [2 marks]	51

Heuristic Search — Theoretical Foundations

This section covers the **theory** that underpins every search-algorithm question in the exam. The year-by-year traces in later sections apply these concepts to specific graphs.

BFS (Uninformed) vs. Best-First Search (Informed)

The fundamental divide in search algorithms is between those that use **no domain knowledge** and those that use **heuristic guidance**.

Criterion	BFS (Uninformed / Blind)	Best-First Search (Informed / Heuristic)
What guides expansion?	Structural position — depth from the root. BFS expands level by level, regardless of where the goal is.	Evaluation function $f(n)$ — incorporates domain knowledge (estimates of cost to the goal). Expands the “most promising” node first.
Fringe data structure	FIFO queue — first-in, first-out. Nodes are processed in the order they are discovered.	Priority queue — nodes are dequeued in order of $f(n)$ (lowest first). The queue is dynamically reordered as new nodes are discovered.
Knowledge used	None — only the problem definition (initial state, goal test, successor function, step costs).	Heuristic function $h(n)$ — estimated cheapest cost from n to the goal. Domain-specific (e.g., straight-line distance, Manhattan distance).
Efficiency	Expensive for large state spaces — $O(b^d)$ time and space. Must explore all nodes at depth $d-1$ before checking depth d .	Potentially much faster — a good heuristic dramatically prunes the search space. But with a poor heuristic, can be as bad as uninformed search.
Completeness	Complete — guaranteed to find a solution if one exists (finite branching factor).	Depends on the variant — Greedy BFS is incomplete (can loop). A* is complete.
Optimality	Optimal only for uniform step costs.	Depends on the variant — Greedy BFS is not optimal. A* is optimal if h is admissible.
Examples	BFS, UCS (expand by $g(n)$), DFS, IDDFS	Greedy Best-First (expand by $h(n)$), A* (expand by $f(n) = g(n) + h(n)$), IDA*, SMA*

Key insight: “Uninformed” does *not* mean “stupid” — it means the algorithm works for **any** problem without modification. “Informed” means the algorithm exploits **domain-specific** knowledge to work more efficiently, but the heuristic must be designed for each problem. An uninformed algorithm needs no setup; an informed algorithm needs a clever heuristic.

Informed vs. Uninformed Search — The Fundamental Divide

Every search algorithm belongs to one of two families, distinguished by a single question:
Does it use domain-specific knowledge beyond the problem definition?

Criterion	Uninformed Search (Blind)	Informed Search (Heuristic)
Knowledge used	Only the problem definition: initial state, goal test, successor function, step costs. No estimate of how close a state is to the goal.	Uses a heuristic function $h(n)$ — domain-specific knowledge estimating the cheapest cost from n to the goal.
How nodes are chosen	By structural criteria — depth (BFS, DFS, IDDFS), path cost so far (UCS), or symmetry (Bidirectional).	By an evaluation function $f(n)$ combining $g(n)$ and $h(n)$ — the algorithm steers toward promising-looking nodes.
Efficiency	Exponential in the worst case. Can explore vast portions of the state space with no guidance.	Potentially far more efficient — a good heuristic prunes large unpromising regions. With a perfect heuristic, A* expands only nodes on the optimal path.
Generality	Universal — works for <i>any</i> problem without modification. No domain knowledge needed. Zero setup cost.	Problem-specific — the heuristic must be designed (or learned) for each problem.
Optimality guarantee	UCS guarantees optimality for any non-negative costs; BFS/IDDFS only for uniform costs.	A* guarantees optimality if h is admissible ($h \leq h^*$). Greedy BFS gives no optimality guarantee.
When to use	Problem is new/unknown; no good heuristic exists; state space is small; or as a fallback.	Domain knowledge is available; an admissible heuristic can be defined; state space is large; optimality matters.
Examples	BFS, DFS, DLS, IDDFS, UCS, Bidirectional Search	Greedy BFS, A*, IDA*, RBFS, SMA*, Beam Search

The continuum — not a hard binary: Some algorithms blur the line.

- UCS uses $g(n)$ — this is “cost-aware” but not “goal-aware.” It knows how much you’ve spent but has no idea how much remains. It is technically uninformed

because $g(n)$ is derived from the problem definition (step costs), not from domain knowledge.

- **Bidirectional Search** uses two simultaneous searches (forward from start, backward from goal). It is **structurally informed** — leveraging the fact that two smaller searches are cheaper than one large one — but uses no heuristic. It is formally uninformed.
- **A*** sits at the intersection: $f(n) = g(n) + h(n)$. The $g(n)$ part is uninformed (known costs); the $h(n)$ part is informed (domain estimate). This is why A* subsumes UCS (when $h(n)=0$) and Greedy (conceptually, when $g(n)$ is ignored).


Progression of Uninformed Search Algorithms

Question: Discuss how uninformed search algorithms can be seen as an evolution of ideas, where each new algorithm is designed to address a limitation of an earlier one. Include the progression among: BFS, DFS, DLS, IDDFS, UCS, and Bidirectional Search.

Uninformed search algorithms did not emerge in isolation — they form a **logical progression**, each motivated by a specific weakness of its predecessor. Understanding this progression turns six seemingly disconnected algorithms into one coherent story.

Act I — The Foundational Trade-off: BFS vs. DFS

The two simplest algorithms represent opposite ends of a fundamental trade-off.



The dilemma: BFS guarantees the answer but needs impossible amounts of memory. DFS uses negligible memory but may never find the answer. The entire subsequent evolution is about **getting the best of both**.

Act II — DFS's First Fix: Depth-Limited Search (DLS)

Motivation	What DLS does	Remaining problem
DFS can descend infinitely deep on an infinite branch, never backtracking to find a shallow solution.	DLS adds a depth cutoff L : DFS stops at depth L and backtracks. This prevents infinite descent — the search is now bounded .	How to choose L? If $L < d$ (solution depth), DLS returns failure even though a solution exists. If $L \gg d$, DLS wastes time. DLS is not complete unless you already know d .

The contribution: DLS showed that adding a depth bound fixes DFS's completeness problem — but only if you know the right bound. The question becomes: *how to find d without knowing it in advance?*

Act III — Finding the Right Depth: Iterative Deepening DFS (IDDFS)

Motivation	What IDDFS does	Remaining problem
DLS fails if $L < d$ (misses solution) or wastes time if $L \gg d$. You don't know d in advance.	IDDFS runs DLS repeatedly with $L = 0, 1, 2, \dots$ until the goal is found. This guarantees finding the shallowest solution — completeness restored .	Non-uniform costs ignored. IDDFS finds the shallowest goal (fewest steps), not the cheapest. If edge costs vary, the shortest path in steps may be expensive. Nodes near the root are re-expanded — overhead factor $\frac{b}{b-1}$.

The contribution: IDDFS solved the BFS-vs-DFS dilemma for uniform-cost problems: $O(bd)$ space (from DFS) + completeness + shallowest-solution optimality (from BFS). The repeated-work overhead is only $\sim 11\%$ for $b=10$ — an acceptable price. **This is the best uninformed algorithm for uniform-cost problems with large state spaces.**

Act IV — Handling Non-Uniform Costs: Uniform-Cost Search (UCS)

Motivation	What UCS does	Remaining problem
BFS and IDDFS assume all steps cost the same . In the real world, actions have different costs (mountain roads cost more than highways). Neither shallowest-path algorithm can find the cheapest path with non-uniform costs.	UCS expands nodes by lowest path cost $g(n)$, not depth. It is optimal for any non-negative edge costs — the first time it expands the goal, the cheapest path is found.	No guidance toward the goal. UCS expands equally in <i>all</i> directions — it knows what you've spent but has no idea how much remains . In a large graph, UCS may explore enormous regions away from the goal.


The contribution: UCS generalised optimality from “shallowest” to “cheapest.” But it exposed a new problem: without knowledge of where the goal is, the algorithm is undirected — it wastes effort exploring in wrong directions. This limitation **directly motivates informed search** — add a heuristic $h(n)$ to estimate remaining distance, and UCS becomes A*.

Act V — Exploiting Symmetry: Bidirectional Search

Motivation	What Bidirectional does	Remaining problem
A single-direction search explores $O(b^d)$ nodes. But a search backward from the goal is equally valid — and two searches meeting in the middle each explore only $O(b^{d/2})$ nodes. This is an exponential reduction : $b^d \rightarrow 2 \cdot b^{d/2}$.	Run two simultaneous searches — one forward from start, one backward from goal(s). The solution is found when the two frontiers intersect. Can use BFS, UCS, or even A* as the underlying search method.	Requires a reversible successor function. The backward search needs a predecessor function — “what states lead to the goal?” This is not always easy to define. Also, checking frontier intersection at every step adds overhead. For some problems (theorem proving), the backward branching factor is much larger.

The contribution: Bidirectional search exploits problem **symmetry** — the fact that searching from the goal backward is just as valid as from the start forward. It is still uninformed, but achieves an exponential speedup through structural cleverness alone. Combined with IDDFS (bidirectional IDDFS) or A* (bidirectional A*), it becomes even more powerful.

The Complete Progression — Visual Summary



↑ Each arrow = a specific limitation being addressed

The Evolutionary Narrative (Exam Summary)

1. **BFS** — Gives us completeness and optimality, but at the impossible price of $O(b^d)$ memory. **Starting point.**
2. **DFS** — Flips the trade-off: $O(bm)$ memory is practical, but the algorithm may never find a solution. **Motivated by BFS's memory explosion.**
3. **DLS** — Fixes DFS's infinite-descent problem with a depth cutoff. But now you must guess the cutoff — guess too low and it reports failure when a solution exists. **Motivated by DFS's incompleteness.**
4. **IDDFS** — Eliminates the guesswork: $L = 0, 1, 2, \dots$ until success. Restores completeness + shallowest-solution optimality while keeping $O(bd)$ space — the best of BFS and DFS. **Motivated by DLS's dependence on unknown L .**
5. **UCS** — Abandons depth for path cost $g(n)$. Finds the cheapest path regardless of non-uniform edge costs. But expands blindly in all directions. **Motivated by IDDFS/BFS's uniform-cost assumption.**
6. **Bidirectional Search** — Attacks exponential complexity directly: two $d/2$ -depth searches replace one d -depth search. Orthogonal improvement — can be layered on top of BFS, UCS, or A*. **Motivated by single-direction search's $O(b^d)$ cost.**

The arc to informed search: The progression above solves structural problems (memory, completeness, cost-sensitivity) but never addresses the **core inefficiency** — none of these algorithms knows where the goal is. The next evolutionary step — **informed search** (Greedy, A*, IDA*) — adds a heuristic $h(n)$ that *estimates*

remaining distance. UCS with $h(n)$ becomes A*; IDDFS with $h(n)$ becomes IDA*. The progression continues, now with domain knowledge.

Heuristic Search — Introduction & the Best-First Family

What Makes a Search “Informed”?

An **informed** (or heuristic) search algorithm uses **problem-specific knowledge** beyond the problem definition to guide the search toward the goal more efficiently. This knowledge is encoded in a **heuristic function** $h(n)$.

Generic Best-First Search Algorithm:

1. Initialise the fringe (priority queue) with the start node.
2. **Loop:** Remove the node with the **lowest evaluation function** $f(n)$ from the fringe.
3. If that node is the goal, return the solution path.
4. Expand the node — generate all successors.
5. For each successor: compute $f(\text{successor})$, insert into the fringe (priority queue).
6. Repeat until the fringe is empty (failure).

The Best-First Search Family: The “family” is unified by the priority-queue fringe — the difference is **what** $f(n)$ measures:

Algorithm	Evaluation $f(n)$	Type	Key property
UCS	$f(n) = g(n)$	Uninformed	Optimal for any non-negative costs.
Greedy BFS	$f(n) = h(n)$	Informed	Fast but not optimal. Uses only heuristic.
A*	$f(n) = g(n) + h(n)$	Informed	Optimal if h is admissible. Best-known informed
IDA*	$f(n) = g(n) + h(n)$	Informed	Like A* but $O(bd)$ space. Uses iterative deepening
RBFS	$f(n) = g(n) + h(n)$	Informed	Recursive Best-First — linear space, slightly more
SMA*	$f(n) = g(n) + h(n)$	Informed	Simplified Memory-Bounded A* — uses all avail

Exam relevance: The CUET exam tests UCS, Greedy BFS, A*, and IDA*/IDDFS — these are the four you must be able to trace step-by-step on a given graph. RBFS and SMA* are mentioned in Russell & Norvig but have never appeared in a CUET past paper (2020–2024).

The Heuristic Function $h(n)$

Definition


$$h(n) = \text{estimated cheapest cost from node } n \text{ to a goal node}$$

Key properties:

- $h(n)$ is **problem-specific** — you design it using domain knowledge. The same generic search code works for any problem, but $h(n)$ encodes what you know about *this* problem.
- $h(n) \geq 0$ for all n , and $h(G) = 0$ for any goal node G .
- $h(n)$ is **not** the action cost — it is an **estimate** of the remaining cost.
- A **perfect heuristic** satisfies $h(n) = h^*(n)$ for all n (where $h^*(n)$ is the true optimal cost). With a perfect heuristic, A* expands only nodes on optimal paths — it does zero wasted work.

Example — Romania Map with Straight-Line Distance Heuristic

The classic illustration from Russell & Norvig (Chapter 3). The problem: drive from **Arad** to **Bucharest** via the Romanian road network. The heuristic $h_{\text{SLD}}(n)$ is the **straight-line (as-the-crow-flies) distance** from city n to Bucharest — the shortest possible distance, ignoring roads.



City n	$h_{\text{SLD}}(n)$ (straight-line km)	True road distance $h^*(n)$
Arad	366	418
Bucharest	0	0
Craiova	160	238
Dobreta	242	360
Fagaras	176	211
Lugoj	244	354
Mehadia	241	353
Oradea	380	531
Pitesti	100	101
Rimnicu	193	193
Sibiu	253	278
Timisoara	329	469
Zerind	374	449

How h_{SLD} works:

- For **any** two points on a map, the straight-line distance is a **lower bound** on the road distance — you cannot drive shorter than a straight line.
- Therefore $h_{\text{SLD}}(n) \leq h^*(n)$ for every city — **the heuristic is admissible.**
- **Check:** Arad: $366 \leq 418$ ✓. Pitesti: $100 \leq 101$ (nearly perfect!). Rimnicu: $193 = 193$ (perfect!). Fagaras: $176 \leq 211$ ✓.
- **No city has $h > h^*$** — the heuristic never overestimates, guaranteeing A* optimality on this problem.
- The heuristic is **not perfect** — Fagaras looks deceptively close (176) but the actual road is 211 (the Carpathian mountains force a winding route). This is exactly why Greedy fails on this problem.

Greedy Best-First Search — Algorithm, Example, Properties

Algorithm

Greedy Best-First Search uses **only the heuristic** as its evaluation function:

$$f(n) = h(n)$$

At every step, expand the fringe node with the **lowest heuristic value** $h(n)$. The fringe is a **min-priority queue sorted by $h(n)$** . Path cost $g(n)$ is **completely ignored**.

Romania Trace — How Greedy Finds a Suboptimal Path

Problem: drive from **Arad** to **Bucharest** using h_{SLD} .

Step	Expanded	$h(n)$	Fringe (sorted by h)
1	Arad	366	S[253], T[329], Z[374] Neighbours: Sibiu(253), Timisoara(329), Zerind(374). Sibiu looks closest.
2	Sibiu	253	F[176], R[193], T[329], A[366], Z[374], O[380] Fagaras has $h=176$ — lowest! Greedy rushes to- ward it.
3	Fagaras	176	B[0], R[193], T[329], A[366], Z[374], O[380] Bucharest has $h=0$ — it is the goal!
4	Bucharest	0	GOAL FOUND

Greedy path: Arad $\xrightarrow{140}$ Sibiu $\xrightarrow{99}$ Fagaras $\xrightarrow{211}$ Bucharest $\Rightarrow$ **Cost = 450**

Optimal path (A* finds this; Greedy misses it): Arad $\xrightarrow{140}$ Sibiu $\xrightarrow{80}$ Rimnicu $\xrightarrow{97}$ Pitesti $\xrightarrow{101}$ Bucharest $\Rightarrow$ **Cost = 418**

Why Greedy failed — the root cause:

At Sibiu, Greedy compared two routes eastward:

- **Fagaras:** $h=176$ — looks just 176 km from Bucharest as the crow flies!
- **Rimnicu:** $h=193$ — looks slightly further.

Greedy chose Fagaras because $176 < 193$, **completely ignoring** that:

- The road from Fagaras to Bucharest is 211 km (mountains force a winding route — h underestimates by 35 km).
- The Rimnicu $\rightarrow$ Pitesti $\rightarrow$ Bucharest route costs $80 + 101 = 181$ km total — h for Rimnicu is only off by 0 (it's perfect!) and Pitesti is off by just 1.

Moral: A low $h(n)$ can be **deceptive**. A node that looks close to the goal may have a very expensive road ahead. Greedy's refusal to track how much it has already spent ($g(n)$) is its fatal flaw.

Properties of Greedy Best-First Search

Property	Status	Explanation
Completeness	× Not complete	Can get stuck in infinite loops — if the heuristic guides it into a cycle, it may never escape. A closed set fixes this but costs $O(b^d)$ memory. On finite graphs with a closed set, it is complete.
Optimality	× Not optimal	Ignores $g(n)$ — will take a cheap-looking path even if the total cost is high. Romania: cost 450 vs. optimal 418.
Time	$O(b^m)$ worst case	Can explore the entire tree with a poor heuristic. With a good heuristic, extremely fast — Romania: only 4 expansions.
Space	$O(b^m)$	Stores all generated nodes in the priority queue. Good heuristic $\Rightarrow$ fewer nodes generated.
Sensitivity	Very high — entirely dependent on h quality.	With perfect $h = h^*$: optimal in $O(d)$ expansions. With $h = 0$ everywhere: degenerates to random selection from the fringe.

Greedy vs. A* — Direct Comparison


Both are **informed search algorithms** using a priority queue. The critical difference is **what they optimise**:

$$\text{Greedy: } f(n) = h(n)$$

$$\text{A*: } f(n) = g(n) + h(n)$$

Criterion	Greedy Best-First	A* Search
Priority	Lowest $h(n)$ only — “Which node <i>looks</i> closest to the goal?”	Lowest $g(n) + h(n)$ — “Which node gives the cheapest estimated total path?”
Completeness	× Not complete (can loop without closed set)	✓ Complete on finite graphs
Optimality	× Not optimal — ignores past cost	✓ Optimal if h is admissible ($h \leq h^*$)
Time	$O(b^m)$ worst case; very fast with good h	$O(b^d)$ worst case; $O(d)$ with perfect h
Space	$O(b^m)$ — stores fringe	$O(b^d)$ — stores fringe (A*’s main practical weakness)
Heuristic req.	None (works with any h)	Must be admissible for optimality; consistent for efficiency on graphs
Romania result	4 expansions, cost 450 (suboptimal)	6 expansions, cost 418 (optimal)
Use when	Speed > optimality; excellent heuristic; large state space	Optimality required; admissible heuristic; state space fits in memory

The core intuition — a simple counterexample where Greedy loses:



G is the goal. S connects directly to G (cost 100). Or $S \rightarrow A \rightarrow G$ (cost $10 + 10 = 20$).

Greedy: At S: $h(A)=5$, $h(G)=0$. Greedy picks G (lowest h), finds path $S \rightarrow G$ at cost **100**. Done. **A*:** At S: $f(A) = 10 + 5 = 15$, $f(G) = 100 + 0 = 100$. A* expands A first, then $A \rightarrow G$ at cost **20**. **Optimal.**

Greedy was fooled — G looks perfect ($h=0$, it *is* the goal!), but reaching it directly costs 100. Path A looks slightly less perfect ($h=5$) but costs only 10 to reach and only 10 more to G. A* accounts for both g and h ; Greedy is blind to g .

Examination 2024 — Q2 (9 Marks)

Problem Data

States: A (start), B, C, D, E, F, H, J, K, G (goal). All roads bidirectional.

Table 1: Road Connections and Step Costs


From	To	Cost
A	B	3
A	C	4
A	D	7
B	E	6
B	F	5
C	F	11
C	H	4
D	H	2
D	J	10
E	G	7
E	H	3
F	G	9
F	J	4
H	G	5
H	K	6
J	K	3
K	G	4

Table 2: Heuristic Estimates $h(n)$ to Goal G

Node	$h(n)$
A	11
B	10
C	7
D	6
E	6
F	6
H	4
J	5
K	3
G	0

Part a(i) — State-Space Graph [1 mark]

Question: Using Table 1, draw the state-space graph clearly labelling all nodes and edge costs.



Note: The edge E–H is curved upward and F–J is curved downward to avoid visual ambiguity where nodes share the same y -coordinate. All edges are **bidirectional** (undirected).

Part a(ii) — Search Algorithm Traces [5 marks]

Question: Show search trees for (1) UCS, (2) Greedy Best-First, (3) A*, (4) IDDFS from A to G. At every step: show node being expanded and fringe contents in order of selection. For IDDFS: show depth limit per iteration. Ties broken alphabetically.

Algorithm 1 — Uniform-Cost Search (UCS)

Key idea: Always expand the node with the **lowest path cost** $g(n)$ from the start. The fringe is a **min-priority queue sorted by** $g(n)$. A node is closed (never re-expanded) once popped. If a *cheaper* path to a fringe node is found, its cost is updated (replaced).

Fringe notation: Node[g] sorted ascending by g .

Step	Node Expanded	$g(n)$	Fringe after expansion (sorted by g)
1	A	0	B[3], C[4], D[7] Neighbours of A: B($g=3$), C($g=4$), D($g=7$).
2	B	3	C[4], D[7], F[8], E[9] A already closed. $B \rightarrow F$: $g=8$; $B \rightarrow E$: $g=9$. Both new.
3	C	4	D[7], F[8], H[8], E[9] $C \rightarrow F$: $g=15 > 8$ already in fringe $\Rightarrow$ keep F[8]. $C \rightarrow H$: $g=8$ new.
4	D	7	F[8], H[8], E[9], J[17] $D \rightarrow H$: $g=9 > 8$ in fringe $\Rightarrow$ keep H[8]. $D \rightarrow J$: $g=17$ new. Tie F,H: F before H (alphabetical).
5	F	8	H[8], E[9], J[12], G[17] $F \rightarrow G$: $g=17$ new. $F \rightarrow J$: $g=12 < 17$ in fringe $\Rightarrow$ update J[12].
6	H	8	E[9], J[12], G[13], K[14] $H \rightarrow E$: $g=11 > 9$ in fringe $\Rightarrow$ keep E[9]. $H \rightarrow G$: $g=13 < 17 \Rightarrow$ update G[13]. $H \rightarrow K$: $g=14$ new.
7	E	9	J[12], G[13], K[14] $E \rightarrow G$: $g=16 > 13$ in fringe $\Rightarrow$ keep G[13].
8	J	12	G[13], K[14] $J \rightarrow K$: $g=15 > 14$ in fringe $\Rightarrow$ keep K[14].
9	G	13	GOAL FOUND

UCS Solution: $A \xrightarrow{4} C \xrightarrow{4} H \xrightarrow{5} G$ **Total cost:** $4 + 4 + 5 = 13$
 Nodes expanded (in order): A, B, C, D, F, H, E, J, G (9 nodes total)
 UCS **guarantees the optimal path** because it always expands the cheapest unexplored node, never re-expands a closed node, and path costs are non-negative.

Algorithm 2 — Greedy Best-First Search

Key idea: Always expand the node with the **lowest heuristic estimate** $h(n)$ to the goal. The fringe is sorted by $h(n)$ (ascending). Path cost $g(n)$ is *completely ignored*. Greedy is fast but not guaranteed to find the optimal path.

Fringe notation: Node[h] sorted ascending by h .

Step	Node Expanded	$h(n)$	Fringe after expansion (sorted by h)
1	A	11	D[6], C[7], B[10] Neighbours B($h=10$), C($h=7$), D($h=6$). Sorted by h : D, C, B.
2	D	6	H[4], J[5], C[7], B[10] D's neighbours excl. A: H($h=4$), J($h=5$).
3	H	4	G[0], K[3], J[5], E[6], C[7], B[10] H's neighbours excl. D: C($h=7$, already in fringe—skip), E($h=6$), G($h=0$), K($h=3$).
4	G	0	GOAL FOUND

Greedy Solution: $A \xrightarrow{7} D \xrightarrow{2} H \xrightarrow{5} G$ **Total cost:** $7 + 2 + 5 = 14$ (not optimal!)

Nodes expanded: A, D, H, G (only 4 nodes — very fast)

Greedy expanded just 4 nodes versus UCS's 9, but found a **suboptimal path** (cost 14 vs optimal 13). It followed the heuristic blindly: D had the lowest h from A, leading it down the D–H–G route instead of the cheaper A–C–H–G route.

Algorithm 3 — A* Search

Key idea: Expand the node with the **lowest evaluation function** $f(n) = g(n) + h(n)$, where $g(n)$ = exact cost from start and $h(n)$ = heuristic estimate to goal. The fringe is sorted by $f(n)$ (ascending). Ties in f broken alphabetically. If a cheaper path to a fringe node is found, update its entry.

Fringe notation: Node[g, h, f] sorted ascending by f .

Step	Node	g	h	f	Fringe after expansion (sorted by f)
1	A	0	11	11	C[4,7, 11], B[3,10, 13], D[7,6, 13] Tie B,D at $f=13$: B first.
2	C	4	7	11	H[8,4, 12], B[3,10, 13], D[7,6, 13], F[15,6, 21] $C \rightarrow H$: $g=8, f=12$ new. $C \rightarrow F$: $g=15, f=21$ new.
3	H	8	4	12	B[3,10, 13], D[7,6, 13], G[13,0, 13], E[11,6, 17], K[14,3, 17], F[15,6, 21] $H \rightarrow D$: $f=16 > 13$, keep. $H \rightarrow G$: $f=13$ new. $H \rightarrow E$: $f=17$ new. $H \rightarrow K$: $f=17$ new. Tie B,D,G at $f=13$: B first .
4	B	3	10	13	D[7,6, 13], G[13,0, 13], F[8,6, 14], E[9,6, 15], K[14,3, 17] $B \rightarrow F$: $g=8, f=14 < 21 \Rightarrow$ update F. $B \rightarrow E$: $g=9, f=15 < 17 \Rightarrow$ update E. Tie D,G at $f=13$: D first .
5	D	7	6	13	G[13,0, 13], F[8,6, 14], E[9,6, 15], K[14,3, 17], J[17,5, 22] $D \rightarrow J$: $g=17, f=22$ new.
6	G	13	0	13	GOAL FOUND

A* Solution: $A \xrightarrow{4} C \xrightarrow{4} H \xrightarrow{5} G$ **Total cost:** $4 + 4 + 5 = 13$ (optimal!)
 Nodes expanded: A, C, H, B, D, G (6 nodes — fewer than UCS's 9)
 A* is both **complete** and **optimal** when the heuristic is *admissible* ($h(n) \leq h^*(n)$, i.e. it never overestimates). It is more efficient than UCS because $h(n)$ guides the search toward promising nodes, pruning unpromising directions early.

Algorithm 4 — Iterative Deepening DFS (IDDFS)

Key idea: Perform DFS with a depth cutoff $L = 0, 1, 2, \dots$, restarting from the root each time, until the goal is found. The “fringe” is a **stack** (LIFO). Children are pushed in *reverse* alphabetical order so they are *popped* (and expanded) in alphabetical order. Nodes on the **current path** (ancestors) are not revisited to avoid infinite loops.

Stack notation: Leftmost entry = top of stack (next to be popped). Depth shown in parentheses.

Iteration $L = 0$:

Step	Action	Stack (top $\leftarrow$ left)
1	Pop A (d=0). $A \neq G$. At depth limit — cannot expand.	[]
<i>No solution at $L = 0$.</i>		

Iteration $L = 1$:

Step	Action	Stack
1	Pop A (d=0). Expand: push D,C,B (B on top).	B(1), C(1), D(1)
2	Pop B (d=1). $B \neq G$. At depth limit.	C(1), D(1)
3	Pop C (d=1). $C \neq G$. At depth limit.	D(1)
4	Pop D (d=1). $D \neq G$. At depth limit.	[]
<i>No solution at $L = 1$.</i>		

Iteration $L = 2$:

Step	Action	Stack
1	Pop A (d=0). Expand: B,C,D.	B(1), C(1), D(1)
2	Pop B (d=1). Expand (excl. A): E,F. Push F,E.	E(2), F(2), C(1), D(1)
3	Pop E (d=2). $E \neq G$. At limit.	F(2), C(1), D(1)
4	Pop F (d=2). At limit.	C(1), D(1)
5	Pop C (d=1). Expand (excl. A): F,H. Push H,F.	F(2), H(2), D(1)
6	Pop F (d=2). At limit.	H(2), D(1)
7	Pop H (d=2). At limit.	D(1)
8	Pop D (d=1). Expand (excl. A): H,J. Push J,H.	H(2), J(2)
9	Pop H (d=2). At limit.	J(2)
10	Pop J (d=2). At limit.	[]
<i>No solution at $L = 2$.</i>		

Iteration $L = 3$:

Step	Action	Stack
1	Pop A (d=0). Expand: B,C,D.	B(1), C(1), D(1)
2	Pop B (d=1). Expand (excl. A): E,F.	E(2), F(2), C(1), D(1)
3	Pop E (d=2). Expand (excl. B): G,H. Push H first then G (G on top).	G(3), H(3), F(2), C(1)
4	Pop G (d=3). $G = \text{GOAL}$.	GOAL FOUND

IDDFS Solution: $A \xrightarrow{3} B \xrightarrow{6} E \xrightarrow{7} G$ **Depth: 3** **Cost:** $3 + 6 + 7 = 16$

IDDFS found the solution at depth-limit $L = 3$ in just 4 node expansions during the final iteration.

Important distinction: IDDFS finds the *shallowest* path (fewest hops), **not** the lowest-cost path. The optimal path $A \rightarrow C \rightarrow H \rightarrow G$ also has depth 3, but alphabetical DFS expansion reaches G via $B \rightarrow E \rightarrow G$ first. For non-uniform step costs, IDDFS is **not optimal**. It is however **complete** — it will always find a solution if one exists.

Why IDDFS over BFS? IDDFS uses only $O(bL)$ space (a single stack), whereas BFS requires $O(b^L)$ space to store all frontier nodes. The repeated work is a constant factor overhead.

Algorithm Comparison Summary:

Algorithm	Path Found	Cost	Nodes Expanded	Optimal?
UCS	$A \rightarrow C \rightarrow H \rightarrow G$	13	9 (all iterations)	Yes
Greedy	$A \rightarrow D \rightarrow H \rightarrow G$	14	4	No
A*	$A \rightarrow C \rightarrow H \rightarrow G$	13	6	Yes
IDDFS	$A \rightarrow B \rightarrow E \rightarrow G$	16	4 (final iter.)	No (non-uniform costs)

Part b — Problem Formulations [3 marks]

Question: Give the initial state, goal test, operators, and path cost function for each problem.

(i) Robot Vacuum Cleaning Problem (Indoor Environment)

State: A pair (l, D) where l is the robot's current room and $D \subseteq \{\text{all rooms}\}$ is the set of dirty rooms. *Example:* (Kitchen, {Bedroom, Lounge}).

Initial state: (start room, {all currently dirty rooms}).

Goal test: $D = \emptyset$ (no dirty rooms remain).

Operators:

1. $\text{Move}(r_i \rightarrow r_j)$: Move from room r_i to adjacent room r_j (changes location; D unchanged).
2. $\text{Suck}()$: Vacuum the current room (removes current room from D if dirty).

Static obstacles (furniture) constrain which rooms are adjacent.

Path cost: Total energy used: each Move costs c_{move} energy units; each Suck costs c_{suck} units. Goal: minimise total energy.

Informed vs. Uninformed: **Informed search (A*) is more suitable.** A strong heuristic $h(n) = |D| \times \bar{d}$ (number of remaining dirty rooms multiplied

by average travel cost between rooms) estimates remaining work without overestimating. With many rooms, uninformed BFS/UCS must explore exponentially many (l, D) state combinations blindly. A* with a good heuristic dramatically prunes this search space.

(ii) Water Jug Problem (5L and 3L Jugs)

State: A pair (x, y) where $x \in \{0, 1, 2, 3, 4, 5\}$ is water (litres) in the 5L jug and $y \in \{0, 1, 2, 3\}$ is water in the 3L jug.

Initial state: $(0, 0)$ — both jugs empty.

Goal test: $x = 4$ (exactly 4 litres in the 5L jug). The 3L jug cannot hold 4L, so only $x = 4$ satisfies the goal.

Operators:

1. **Fill 5L:** $(x, y) \rightarrow (5, y)$
2. **Fill 3L:** $(x, y) \rightarrow (x, 3)$
3. **Empty 5L:** $(x, y) \rightarrow (0, y)$
4. **Empty 3L:** $(x, y) \rightarrow (x, 0)$
5. **Pour 5L $\rightarrow$ 3L:** $(x, y) \rightarrow (x-d, y+d)$ where $d = \min(x, 3-y)$
6. **Pour 3L $\rightarrow$ 5L:** $(x, y) \rightarrow (x+d, y-d)$ where $d = \min(y, 5-x)$

Path cost: Number of actions performed (1 per action), minimised to find the shortest sequence.

State space: **FINITE.** The 5L jug holds integer amounts $\{0, 1, 2, 3, 4, 5\}$ (6 values) and the 3L jug holds $\{0, 1, 2, 3\}$ (4 values), giving at most $6 \times 4 = 24$ states. Not all are reachable, but the space is strictly bounded.

One valid solution path (6 steps):

$$(0, 0) \xrightarrow{\text{Fill 5L}} (5, 0) \xrightarrow{\text{Pour 5} \rightarrow \text{3}} (2, 3) \xrightarrow{\text{Empty 3L}} (2, 0) \xrightarrow{\text{Pour 5} \rightarrow \text{3}} (0, 2) \xrightarrow{\text{Fill 5L}} (5, 2) \xrightarrow{\text{Pour 5} \rightarrow \text{3}} (4, 2)$$

Examination 2023 — Q2 (9 Marks)

Problem Data

Table 1: State Transitions and Step

Costs		
State	Next	Cost
A	B	4
A	C	1
B	D	3
B	E	8
C	C	0
C	D	2
C	F	6
D	C	2
D	E	4
E	G	2
F	G	8

Table 2: Heuristic Estimates $h(n)$

Node	$h(n)$
A	8
B	8
C	6
D	5
E	1
F	4
G	0

Important: This is a **directed** graph — edges go one way only (State $\rightarrow$ Next). Note the self-loop $C \rightarrow C$ with cost 0 — this is a degenerate edge and should be ignored during search (it does not change the state or advance toward the goal). Also note the edge $D \rightarrow C(2)$: returning from D to C is allowed since the graph is directed — C is reachable from both A and D. No reverse edges from $B \rightarrow A$, $C \rightarrow A$, $D \rightarrow B$, etc. exist unless listed in the table above.

Part a — Problem Formulations [2 marks]

Question: Give the initial state, goal test, operators, and path cost for each problem.

(i) Travelling Salesman Problem (TSP)

- State:** A pair (ℓ, V) where ℓ is the current city and $V \subseteq \{\text{all } N \text{ cities}\}$ is the set of *unvisited* cities. *Example midway:* (Dhaka, {Chittagong, Sylhet}) means we are in Dhaka and still need to visit Chittagong and Sylhet.
- Initial state:** (start city, {all $N-1$ other cities}).
- Goal test:** $\ell = \text{start city}$ **and** $V = \emptyset$ (returned to start after visiting all cities exactly once).
- Operators:** Travel($\text{city}_i \rightarrow \text{city}_j$) where $\text{city}_j \notin V$ or ($V = \emptyset$ and $\text{city}_j = \text{start}$). Each move removes city_j from V .
- Path cost:** Sum of distances (or costs) of all travelled roads. Goal: find the tour with **minimum total cost**.

Complexity note: With N cities there are $(N-1)!/2$ distinct tours — factorial growth. TSP is NP-hard, so exhaustive search is infeasible for $N > 20$. Heuristic search (A^*) with an admissible heuristic (e.g. minimum spanning tree of remaining cities) can find optimal or near-optimal tours for moderate N .

(ii) Missionaries & Cannibals

State: A triple (M_L, C_L, B) where M_L = missionaries on left bank (0–3), C_L = cannibals on left bank (0–3), $B \in \{L, R\}$ = boat position. The right bank has $3-M_L$ missionaries and $3-C_L$ cannibals.

Initial state: $(3, 3, L)$ — all six people and the boat on the left bank.

Goal test: $(0, 0, R)$ — all six people and the boat on the right bank.


Operators: Move k missionaries and j cannibals across the river **by boat**, where $1 \leq k+j \leq 2$ (boat capacity 2), and the resulting state must be **legal**: on both banks, $M \geq C$ or $M = 0$ (missionaries never outnumbered by cannibals on either side).

Possible boat loads: $(1M, 0C)$, $(2M, 0C)$, $(0M, 1C)$, $(0M, 2C)$, $(1M, 1C)$.

Path cost: Number of boat trips (1 per crossing). Goal: find a sequence with **minimum trips**.

State space size: Without constraints: $4 \times 4 \times 2 = 32$ states. After legality filtering, only about 16–20 are valid. The problem has a known shortest solution of **11 crossings**.

Part b(i) — State-Space Diagram [0.5 mark]



Notes on the graph structure:

- $C \rightarrow C(0)$ is a zero-cost self-loop — degenerate; ignored in search.
- $D \rightarrow C(2)$ creates a directed 2-cycle $C \rightleftharpoons D$ ($C \rightarrow D$ costs 2, $D \rightarrow C$ costs 2).
- B has two outgoing edges: $B \rightarrow D(3)$ and $B \rightarrow E(8)$. $B \rightarrow E$ is an expensive direct path.
- The cheap route to the goal passes through $D \rightarrow E(4) \rightarrow G(2)$.
- From A, the best first step is C (cost 1), not B (cost 4).

Part b(ii) — Search Algorithm Traces [4.5 marks]

Question: Trace UCS, Greedy search, and A* from A (initial) to G (goal). At each step: show the node being expanded and the fringe contents. Ties broken alphabetically.

Algorithm 1 — Uniform-Cost Search (UCS)

Key idea: Expand the node with the **lowest path cost** $g(n)$. Fringe is a min-priority queue sorted by $g(n)$. Once a node is expanded (closed), it is never re-expanded. If a cheaper path to a fringe node is found, update its cost.

Fringe notation: Node[g] sorted ascending by g . **Closed set:** {nodes already expanded}.

Step	Node	g	Fringe after expansion (sorted by g)
1	A	0	C[1], B[4] $A \rightarrow C$: $g=1$. $A \rightarrow B$: $g=4$. Both new.
2	C	1	D[3], B[4], F[7] $C \rightarrow C$: self-loop, ignore. $C \rightarrow D$: $g=3$ new. $C \rightarrow F$: $g=7$ new.
3	D	3	B[4], E[7], F[7] $D \rightarrow C$: $g=5$, C already closed $\Rightarrow$ skip. $D \rightarrow E$: $g=7$ new.
4	B	4	E[7], F[7] $B \rightarrow D$: $g=7$, D already closed $\Rightarrow$ skip. $B \rightarrow E$: $g=12 > E[7]$ in fringe $\Rightarrow$ keep E[7].
5	E	7	F[7], G[9] Tie E[7], F[7]: E before F (alphabetical). $E \rightarrow G$: $g=9$ new.
6	F	7	G[9] $F \rightarrow G$: $g=15 > G[9]$ in fringe $\Rightarrow$ keep G[9].
7	G	9	GOAL FOUND

UCS Solution: $A \xrightarrow{1} C \xrightarrow{2} D \xrightarrow{4} E \xrightarrow{2} G$ **Total cost:** $1 + 2 + 4 + 2 = 9$
 Nodes expanded (in order): A, C, D, B, E, F, G (7 nodes total)
 UCS guarantees the optimal path because it explores in order of increasing $g(n)$ and step costs are non-negative.

Algorithm 2 — Greedy Best-First Search

Key idea: Expand the node with the **lowest heuristic** $h(n)$. Fringe sorted by h ascending. Path cost $g(n)$ is *completely ignored*. This is fast but not guaranteed optimal.

Fringe notation: Node[h] sorted ascending by h .

Step	Node	h	Fringe after expansion (sorted by h)
1	A	8	C[6], B[8] $A \rightarrow C$: $h=6$. $A \rightarrow B$: $h=8$.
2	C	6	F[4], D[5], B[8] $C \rightarrow F$: $h=4$. $C \rightarrow D$: $h=5$. $C \rightarrow C$ ignored.
3	F	4	G[0], D[5], B[8] $F \rightarrow G$: $h=0$. Goal in fringe with lowest h .
4	G	0	GOAL FOUND

Greedy Solution: $A \xrightarrow{1} C \xrightarrow{6} F \xrightarrow{8} G$ **Total cost:** $1 + 6 + 8 = 15$ (suboptimal!)
 Nodes expanded: A, C, F, G (only 4 nodes — extremely fast)
 Greedy blindly chased the heuristic: from C, F has $h=4$ (lowest), so it went $C \rightarrow F$, but $F \rightarrow G$ costs 8 — the most expensive path to the goal. It completely missed the cheap $C \rightarrow D \rightarrow E \rightarrow G$ route (cost 9) because D's $h=5$ looked slightly worse than F's $h=4$.

Algorithm 3 — A* Search

Key idea: Expand the node with the **lowest** $f(n) = g(n) + h(n)$. Fringe sorted by f ascending. Ties in f broken alphabetically. Closed set prevents re-expansion.

Fringe notation: Node[g, h, f] sorted ascending by f .

Step	Node	g	h	f	Fringe after expansion (sorted by f)
1	A	0	8	8	C[1,6, 7], B[4,8, 12] $A \rightarrow C$: $f=1+6=7$. $A \rightarrow B$: $f=4+8=12$.
2	C	1	6	7	D[3,5, 8], F[7,4, 11], B[4,8, 12] $C \rightarrow D$: $g=3, f=8$. $C \rightarrow F$: $g=7, f=11$. $C \rightarrow C$ ignored.
3	D	3	5	8	E[7,1, 8], F[7,4, 11], B[4,8, 12] $D \rightarrow C$: C closed $\Rightarrow$ skip. $D \rightarrow E$: $g=7, f=8$. Tie D($f=8$) and E($f=8$): D already expanded, E is in fringe.
4	E	7	1	8	G[9,0, 9], F[7,4, 11], B[4,8, 12] $E \rightarrow G$: $g=9, f=9$.
5	G	9	0	9	GOAL FOUND

A* Solution: $A \xrightarrow{1} C \xrightarrow{2} D \xrightarrow{4} E \xrightarrow{2} G$ **Total cost:** $1 + 2 + 4 + 2 = 9$ (optimal!)

Nodes expanded: A, C, D, E, G (only 5 nodes — fewer than UCS's 7)

A* found the optimal path and expanded **only 5 nodes** compared to UCS's 7, because the heuristic $h(n)$ guided it away from the expensive B and F branches. Note that **B was never expanded** — its $f=12$ was always worse than the alternatives ahead of it in the fringe.

Algorithm Comparison for 2023 Graph:

Algorithm	Path Found	Cost	Nodes Expanded	Optimal?
UCS	$A \rightarrow C \rightarrow D \rightarrow E \rightarrow G$	9	7	Yes
Greedy	$A \rightarrow C \rightarrow F \rightarrow G$	15	4	No
A*	$A \rightarrow C \rightarrow D \rightarrow E \rightarrow G$	9	5	Yes

Why does A* outperform UCS here? The heuristic $h(n)$ provides a useful lower bound. After expanding A and C, A* sees that D has $f=8$ (promising — $h(D)=5$ is close to the true remaining cost of 6) while B has $f=12$ (unpromising — $h(B)=8$ overestimates its potential). A* expands D next and finds the optimal path without ever touching B. UCS, lacking a heuristic, expands B unnecessarily at step 4 before reaching the goal.

Admissibility check: Is $h(n)$ admissible? Let's verify:

- $h(A)=8 \leq h^*(A)=9$ ✓
- $h(C)=6 \leq h^*(C)=8$ ✓
- $h(D)=5 \leq h^*(D)=6$ ✓
- $h(E)=1 \leq h^*(E)=2$ ✓
- $h(B)=8 \leq h^*(B)=5$? **No!** $h^*(B) = \min(3+4+2, 8+2) = 9$. Wait — $B \rightarrow D \rightarrow E \rightarrow G = 3+4+2 = 9$, so $h(B)=8 < 9$ ✓. And $B \rightarrow E \rightarrow G = 8+2 = 10$. So $h^*(B)=9$, $h(B)=8$ is admissible.

- $h(F)=4 \leq h^*(F)=8$ ✓

Yes, the heuristic is admissible — $h(n) \leq h^*(n)$ for all nodes. This is why A* found the optimal solution.

Part c — IDDFS: Synthesising Benefits of BFS and DFS [2 marks]


Question: Illustrate with an appropriate example how IDDFS synthesises the benefits of both BFS and DFS — achieving optimality, completeness, and linear space complexity. Nevertheless, time complexity deteriorates in comparison to both BFS and DFS.

The Core Idea

IDDFS (Iterative Deepening Depth-First Search) runs DFS repeatedly with an increasing depth cutoff: first $L=0$, then $L=1$, $L=2$, $\dots$, until the goal is found. At each iteration, the search is a standard depth-limited DFS from the root.

Illustrative Example

Consider the following binary tree where G (shaded) is the shallowest goal at depth 3:



IDDFS Execution:

L	Nodes visited (DFS order)	Outcome
0	A	$A \neq G$. No solution at depth 0.
1	A, B, C	Neither B nor C is G. No solution at depth 1.
2	A, B, D, E, C, F, I	No goal at depth 2.
3	A, B, D, G ← GOAL FOUND	Found! Shallowest goal at depth 3.

How IDDFS Inherits Each Algorithm's Strengths

Property	BFS	DFS	IDDFS
Completeness	✓ Complete (finds solution if one exists)	× Incomplete (may loop forever on infinite branches)	✓ Complete — like Depth-limited DFS and avoids infinite loops because it backtracks at the cutoff
Optimality	✓ Optimal for uniform step costs (finds shallowest solution)	× Not optimal (may find a deep solution first)	✓ Optimal for uniform costs — like BFS. Because it explores all nodes at depth $d-1$ before any at depth d , the first solution found is the shallowest
Space	× $O(b^d)$ — stores entire frontier. For $b=10, d=10$, that's 10 billion nodes.	✓ $O(bm)$ — stores only one path + siblings.	✓ $O(bd)$ — like DFS. Only the current path (depth d) plus $b-1$ siblings at each level are in memory. $b=10, d=10$: ≈ 100 nodes
Time	$O(b^d)$ — visits each node once	$O(b^m)$ — may go very deep, but each node visited once	× $O(b^d)$ — worse constant factor . Nodes near the root are revisited many times. Overhead factor is large (e.g. for $b=10$, about 10 times BFS).

Why Time Complexity Deteriorates

The key weakness of IDDFS is **repeated work**. Consider the root node A: it is expanded in *every* iteration ($L=0, 1, 2, 3$) — 4 times. Node B is expanded at $L=2$ and $L=3$ — twice. Only the deepest nodes (depth L) are expanded just once.

However, in a tree with branching factor b , the **number of nodes at the deepest level dominates**:

$$\begin{aligned}
 \text{Total expansions} &= (d+1) \cdot 1 + d \cdot b + (d-1) \cdot b^2 + \dots + 1 \cdot b^d \\
 &= b^d \left(1 + \frac{1}{b} + \frac{1}{b^2} + \dots \right) \approx b^d \cdot \frac{b}{b-1}
 \end{aligned}$$

For $b = 10$, the overhead is only $\frac{10}{9} \approx 1.11\times$ — a small price for the enormous space savings ($O(bd)$ vs $O(b^d)$). This is why IDDFS is the **preferred uninformed search strategy** when the state space is large and step costs are uniform.

When IDDFS Is NOT Optimal

IDDFS finds the **shallowest** goal (fewest edges), not the lowest-cost goal. With non-uniform step costs, a deeper path with lower total cost may exist. In such cases, UCS or A* is required. For the 2024 exam graph, IDDFS found $A \rightarrow B \rightarrow E \rightarrow G$ (cost 16, 3 hops) while UCS found $A \rightarrow C \rightarrow H \rightarrow G$ (cost 13, 3 hops) — both at the same depth, but UCS found the cheaper one because it considers edge costs, not just depth.

Exam summary — IDDFS:

- **Completeness:** Inherited from BFS — guaranteed to find a solution if one exists.
- **Optimality:** Inherited from BFS — optimal for *uniform* step costs (finds shallowest solution).
- **Space:** Inherited from DFS — $O(bd)$ linear space, practical for large search spaces.
- **Time:** Worse than both BFS and DFS individually — nodes near the root are re-expanded each iteration. However, the overhead is only a constant factor ($\approx \frac{b}{b-1}$), acceptable given the space savings.

Examination 2022

Note: Search questions in 2022 are **scattered across two sections**:

- **Section A, Q.A-4(a):** DFS vs BFS theory (2 marks)
- **Section B, Q.B-4(b):** Graph trace — UCS, Greedy, A* (3 marks) — *identical search space to 2023*

Section A — Q.A-4(a): DFS vs BFS [2 marks]

Question: What is an AI Technique? Write down the advantages and disadvantages of Depth-First Search and Breadth-First Search.

AI Technique

An **AI Technique** is a method that exploits knowledge, represented in a form that:

1. **Captures generalisations** — groups situations sharing important properties together, avoiding separate representations for every individual case.
2. **Can be understood** by the people providing it — knowledge must be inspectable and modifiable.
3. **Can be easily modified** to correct errors and reflect new information.
4. **Can be used in many situations** even if not perfectly complete or accurate.
5. **Reduces the search space** by narrowing the range of possibilities that must be considered.

In short: an AI technique is a **knowledge-driven method for controlling search** so that solutions are found efficiently without exhaustive enumeration of all possibilities. DFS and BFS are two fundamental *uninformed* (blind) search techniques — they use no domain-specific knowledge, only the structure of the search tree.

Depth-First Search (DFS)

Advantages	Disadvantages
1. Low memory: Stores only the current path (b siblings per level), giving $O(bm)$ space. For deep trees, this is the decisive advantage.	1. Not complete: Fails on infinite-depth trees. DFS may follow an infinite branch forever and never find a solution elsewhere.
2. Simple implementation: Easy to code recursively or with a stack (LIFO).	2. Not optimal: The first solution found may be very deep in the tree. DFS has no mechanism to prefer shorter paths.
3. Fast to a solution if lucky: If the goal is deep and branching is low, DFS may find it with very few node expansions.	3. Susceptible to loops: Without a closed/visited set, DFS can cycle infinitely.
4. Backtracking-friendly: Natural fit for constraint satisfaction and puzzle-solving where partial solutions are built and abandoned.	4. May miss shallow goals: DFS goes deep first — a goal at depth 2 may be found only after exhausting paths to depth 50.

Breadth-First Search (BFS)

Advantages	Disadvantages
1. Complete: Guaranteed to find a solution if one exists (provided b is finite).	1. Exponential memory: Must store the entire frontier — $O(b^d)$ space. For $b=10$, $d=12$: 1 trillion nodes in memory. This makes BFS infeasible for most real problems.
2. Optimal (uniform cost): Finds the shallowest solution. If step costs are all equal, this is the optimal cost.	2. Exponential time: $O(b^d)$ — same as the space requirement. Practical only for shallow trees.
3. Systematic: Processes the entire tree level by level — no path is ever skipped or overlooked.	3. Slow to find deep goals: Must explore all nodes at depth $d-1$ before checking any at depth d . If the goal is at depth 10, all 9 shallow levels must be fully expanded first.
4. Good for shallow goals: If the goal is near the root, BFS finds it quickly.	4. Queue implementation overhead: Maintaining the frontier as a FIFO queue with $O(1)$ enqueue/dequeue while also tracking visited states adds complexity.

When to Use Which

Criterion	Use DFS	Use BFS
Memory is limited	Yes — $O(bm)$ space	No — $O(b^d)$ space
Goal is known to be deep	Yes — goes deep fast	No — explores all shallow levels first
Goal is known to be shallow	No — may waste time deep	Yes — finds it early
Need shortest path	No — not optimal	Yes (uniform costs)
Infinite / very deep tree	No — may not terminate	Yes — complete

The IDDFS resolution: Iterative Deepening DFS combines the space advantage of DFS ($O(bd)$ memory) with the completeness and optimality of BFS. For most practical uninformed search problems where memory is the bottleneck, **IDDFS is the preferred choice**. It was tested in 2024 Q2, 2023 Q2(c), 2021 Q3(b), and 2020

Q3(b) — a core exam pattern.

Section B — Q.B-4(b): Graph Trace [3 marks]

Question: Suppose you have the following search space (State, Next, Cost as given in Table 1, $h(n)$ as given in Table 2). Assume initial state A and goal state G. Show how each of the following search strategies would create a search tree to find a path from A to G: **Uniform cost, Greedy search, and A***. At each step show which node is being expanded and the content of fringe.

The **2022 search space is identical to the 2023 search space** (same states, same edges, same $h(n)$ values). The full step-by-step traces for UCS, Greedy, and A* on this graph are given in **Section 2.3** (Examination 2023, Part b(ii)). Below is the summary.

Algorithm	Path	Cost	Nodes Expanded
UCS	$A \rightarrow C \rightarrow D \rightarrow E \rightarrow G$	9	7 (A, C, D, B, E, F, G)
Greedy	$A \rightarrow C \rightarrow F \rightarrow G$	15	4 (A, C, F, G)
A*	$A \rightarrow C \rightarrow D \rightarrow E \rightarrow G$	9	5 (A, C, D, E, G)

Key observation — same graph appears 3 times: The identical search space ($A \rightarrow B(4)$, $A \rightarrow C(1)$, $B \rightarrow D(3)$, $B \rightarrow E(8)$, $C \rightarrow D(2)$, $C \rightarrow F(6)$, $D \rightarrow C(2)$, $D \rightarrow E(4)$, $E \rightarrow G(2)$, $F \rightarrow G(8)$) with the same $h(n)$ values is tested in **2020 Q2, 2022 B-4(b), and 2023 Q2**. Mastering the traces on this one graph gives you answers for 3 separate exam years. The cheapest path is always $A \rightarrow C \rightarrow D \rightarrow E \rightarrow G$ costing 9.

Examination 2021

Note: Search questions in 2021 appear across two Q-numbers:

- **Q2 (9.75 marks):** Full search question — problem formulation + graph traces (DFID, Greedy, A*)
- **Q3(b) (1.75 marks):** DFID vs IDA* — comparative theory

Q2(a) — State Space Search Problem [0.75 mark]

Question: What is the state space search problem?

A **state space search problem** is a formal representation of a problem as a directed graph where:

- **States** are vertices — each represents a possible configuration of the world.
- **Actions** (or operators) are directed edges — transitions from one state to another.
- A **solution** is a sequence of actions (a path) from an initial state to a goal state.
- The search algorithm **systematically explores** the state space to find this path.

Formally, a state space search problem is defined by four components:

1. **Initial state** — the starting configuration.
2. **Goal test** — a function that determines whether a given state is a goal.
3. **Successor function** (operators) — given a state, returns all (action, next_state) pairs.
4. **Path cost function** — assigns a numeric cost to each path (typically additive).

The **working principle** of a basic search algorithm:

1. Initialise the fringe with the start state.
2. Loop: remove a state from the fringe; if it satisfies the goal test, return the solution path; otherwise, expand it (generate successors) and add them to the fringe.
3. Repeat until the fringe is empty (failure) or goal is found (success).

Q2(b) — Problem Formulations [3 marks]

Question: Give the initial state, goal test, successor function, cost function, and working principle of a basic search algorithm for each problem.

(i) Map Coloring Problem

Problem: Using only 4 colors, color a planar map such that no two adjacent regions share the same color.

State: A partial assignment of colors $\{R, G, B, Y\}$ to a subset of regions $\{r_1, r_2, \dots, r_N\}$. Unassigned regions are marked as $\emptyset$. *Example:* $\{r_1=R, r_2=G, r_3=\emptyset, \dots\}$.

- Initial state:** All N regions unassigned: $\{r_i = \emptyset \mid i = 1, \dots, N\}$.
- Goal test:** All regions are assigned a color, and for every pair of adjacent regions (r_i, r_j) , $color(r_i) \neq color(r_j)$.
- Successor function:** Pick an unassigned region r_i and assign it any of the 4 colors $\{R, G, B, Y\}$ that is **not** already used by any of r_i 's adjacent (already-colored) regions. If no color is valid, backtrack.
- Path cost:** Typically 1 per assignment (uniform cost). Since all solutions assign exactly N colors, any valid solution has cost N . The goal is to find *any* valid coloring, not necessarily optimal.
- Working principle:** This is a **constraint satisfaction problem (CSP)**. Backtracking DFS is the natural fit: assign a color to the next variable, check constraint consistency, backtrack on conflict. DFS is preferred over BFS here because all solutions are at the same depth N , and DFS's low memory allows backtracking through a deep search tree.

(ii) Monkey & Bananas Problem

Problem: A 3-ft monkey is in a room with an 8-ft ceiling. Bananas hang from the ceiling. Two 3-ft crates are available (stackable, movable, climbable). The monkey must get the bananas.

- State:** A 4-tuple $(monkey_pos, crate1_pos, crate2_pos, has_bananas)$ where positions are discrete locations in the room and $has_bananas \in \{\text{true}, \text{false}\}$. If the monkey is on a crate, its effective height increases by the crate height (3 ft per crate; stacking 2 crates gives 6 ft height). The monkey can reach the bananas only if total height ≥ 8 ft.
- Initial state:** $(floor_position, initial_crate1_pos, initial_crate2_pos, false)$ — monkey on floor, crates at their initial locations, no bananas.
- Goal test:** $has_bananas = \text{true}$.
- Successor function (operators):**
1. **Walk**($x \rightarrow y$): Monkey moves from position x to y on the floor (not on a crate).
 2. **Push**(crate, $x \rightarrow y$): Monkey pushes a crate from x to y (must be on floor, adjacent to crate).
 3. **Climb**(crate): Monkey climbs onto a crate (must be adjacent to it on the floor).
 4. **ClimbDown**(): Monkey climbs down from crate to floor.
 5. **Stack**(crate1, crate2): Place crate1 on top of crate2 (monkey must be on floor, crates adjacent).
 6. **Grasp**(): Monkey grabs bananas (must be within reach — total height ≥ 8 ft).
- Path cost:** Number of actions (1 per action). Goal: find the **shortest** sequence to get the bananas.

Key insight: The monkey needs height ≥ 8 ft. One crate gives $3+3=6$ ft (not enough). Two crates stacked give $3+3+3=9$ ft (enough). So the solution requires: walk to crate1, push crate1 under bananas, walk to crate2, push crate2 to same spot, stack crate2 on crate1, climb, grasp.

Q2(c) — Graph Traces [5 marks]

Problem Data


Table 1: Directed Edges and Costs

From	To	Cost
S	A	2
S	B	5
A	C	4
A	D	3
B	E	6
C	G	5
B	G	4

Table 2: Heuristic Estimates $h(n)$

Node	$h(n)$
S	9
A	9
B	4
C	5
D	∞
E	∞
G	0

Graph reconstruction: The exact edge costs from the 2021 past paper Figure 1 are reconstructed from the given $h(n)$ values. The heuristic is **exact** (admissible and consistent): $h(n) = h^*(n)$ for every node. D and E are dead ends ($h=\infty$ means no path to G exists from these nodes).



Question: Assume start=S, goal=G. Show search trees for **Depth First Iterative Deepening (DFID)**, **Greedy**, and **A*** — at each step show node expanded and the fringe. Also report the eventual algorithm and solution cost.

Algorithm 1 — Depth-First Iterative Deepening (DFID)

Key idea: Perform depth-limited DFS with increasing cutoff $L = 0, 1, 2, \dots$ until goal is found. The fringe is a **stack** (LIFO). At each iteration, restart from S. Nodes are pushed in **reverse alphabetical** order so they are popped in alphabetical order.

Iteration $L = 0$:

Step	Action	Stack
1	Pop S($d=0$). $S \neq G$. At depth limit.	[]

No solution at $L = 0$.

Iteration $L = 1$:

Step	Action	Stack
1	Pop S($d=0$). Expand: push B, then A (A on top).	A(1), B(1)
2	Pop A($d=1$). $A \neq G$. At limit.	B(1)
3	Pop B($d=1$). $B \neq G$. At limit.	[]

No solution at $L = 1$.

Iteration $L = 2$:

Step	Action	Stack
1	Pop S($d=0$). Push B, A (A on top).	A(1), B(1)
2	Pop A($d=1$). Expand (excl. S): C, D. Push D, C (C on top).	C(2), D(2), B(1)
3	Pop C($d=2$). Expand (excl. A): G. Push G. But C is at depth $d=2$, so G would be at depth $3 > L=2$. Cannot push! C at limit. $C \neq G$.	D(2), B(1)
4	Pop D($d=2$). D at limit.	B(1)
5	Pop B($d=1$). Expand (excl. S): E, G. Push G, E (E on top).	E(2), G(2)
6	Pop E($d=2$). E at limit.	G(2)
7	Pop G($d=2$). G = GOAL!	GOAL FOUND

DFID Solution: $S \xrightarrow{5} B \xrightarrow{4} G$ **Depth: 2** **Cost:** $5 + 4 = 9$

DFID finds the **shallowest** goal (depth 2, via $S \rightarrow B \rightarrow G$) at iteration $L = 2$. DFID is complete and optimal for uniform step costs, but here with non-uniform edge costs, it finds the shallowest goal rather than necessarily the cheapest — however, in this graph, $S \rightarrow B \rightarrow G$ happens to also be the optimal-cost path.

Algorithm 2 — Greedy Best-First Search

Fringe notation: Node[h] sorted ascending by h .

Step	Node	h	Fringe after expansion (sorted by h)
1	S	9	B[4], A[9] $S \rightarrow B$: $h=4$. $S \rightarrow A$: $h=9$. B has lowest h , promising.
2	B	4	G[0], E[∞], A[9] $B \rightarrow G$: $h=0$ (goal!). $B \rightarrow E$: $h=\infty$ (dead end, goes to back).
3	G	0	GOAL FOUND

Greedy Solution: $S \xrightarrow{5} B \xrightarrow{4} G$ **Cost:** $5 + 4 = 9$ (optimal — by luck!)

Nodes expanded: S, B, G (only 3 nodes)

Greedy found the optimal path because the heuristic h strongly prefers B ($h=4$) over A ($h=9$), and B directly connects to G. Greedy never even explored the A-branch, saving work.

Algorithm 3 — A* Search

Fringe notation: Node[g, h, f] sorted ascending by f .

Step	Node	g	h	f	Fringe after expansion (sorted by f)
1	S	0	9	9	B[5,4,9], A[2,9,11] $S \rightarrow B$: $f=5+4=9$. $S \rightarrow A$: $f=2+9=11$.
2	B	5	4	9	G[9,0,9], A[2,9,11], E[11, ∞ , ∞] $B \rightarrow G$: $f=9$. $B \rightarrow E$: $f=\infty$, deprioritised. Tie G[9],B[9]: B already expanded, G next.
3	G	9	0	9	GOAL FOUND

A* Solution: $S \xrightarrow{5} B \xrightarrow{4} G$ **Cost:** $5 + 4 = 9$ (optimal!)

Nodes expanded: S, B, G (only 3 nodes — same as Greedy!)

A* mirrors Greedy in this graph because the heuristic is **exact** ($h(n) = h^*(n)$ for all n), meaning $f(n) = g(n) + h^*(n)$ — the exact total cost to the goal through n . When the heuristic is exact, A* expands only nodes on optimal paths. A* never expanded A ($f=11 > 9$) — it was correctly identified as lying on a suboptimal path.

Algorithm Comparison for 2021 Graph:

Algorithm	Path Found	Cost	Nodes Expanded	Optimal?
DFID	$S \rightarrow B \rightarrow G$	9	7 (over all iterations)	Yes (this graph)
Greedy	$S \rightarrow B \rightarrow G$	9	3	Yes (by luck)
A*	$S \rightarrow B \rightarrow G$	9	3	Yes

Why Greedy found the optimal path here: The heuristic is **exact** — $h(n)$ equals the true optimal cost to G from every node. Since $h(B)=4$ ($B \rightarrow G$ costs 4) and $h(A)=9$ ($A \rightarrow C \rightarrow G$ costs 9), Greedy correctly identifies B as the better first step. However, **this is not guaranteed in general**. With a less accurate heuristic, Greedy could be misled. The 2023 graph demonstrated exactly this: $h(F)=4$ looked slightly better than $h(D)=5$, but $F \rightarrow G$ cost 8 while $D \rightarrow E \rightarrow G$ cost 6, making the D-path cheaper overall.

Q3(b) — DFID vs IDA* [1.75 marks]

Question: Differentiate between DFID and IDA* algorithms. Demonstrate how DFID incorporates the advantages of BFS and DFS with an example.


DFID (Depth-First Iterative Deepening)

DFID runs depth-limited DFS with cutoff $L = 0, 1, 2, \dots$ until the goal is found. The cutoff is purely **structural** — it limits the number of edges (depth) from the root.

How DFID combines BFS and DFS advantages:

Property	BFS	DFS	DFID (inherits)
Completeness	✓ Complete	× May loop forever	✓ Complete — BFS. Each iteration bounded depth; branches cannot t
Optimality	✓ Optimal (uniform cost)	× Not optimal	✓ Optimal for uniform costs — BFS. Finds shortest solution first.
Space	× $O(b^d)$	✓ $O(bm)$	✓ $O(bd)$ — from BFS. Only stores one p

Example — Binary Tree (branching factor $b = 2$):



DFID execution on this tree:

- $L=0$: Visit S. Not G.
- $L=1$: Visit S, A, B. Not G.
- $L=2$: Visit S, A, C, **G** ← Found at depth 2.

BFS would: Store all nodes at depth 2 (4 nodes) in the frontier — $O(b^d)$ memory. **DFS would:** Store only the path $S \rightarrow A \rightarrow C$ (3 nodes) — $O(bd)$ memory. **DFID:** Also stores only the path $S \rightarrow A \rightarrow C$ (3 nodes) — $O(bd)$ memory, but with the completeness guarantee of BFS.

IDA* (Iterative Deepening A*)

IDA* is the **informed** counterpart of DFID. Instead of a depth cutoff, it uses an **f -cost cutoff**:

- Each iteration performs DFS, pruning any node where $f(n) = g(n) + h(n) > cutoff$.
- The initial cutoff is $f(start) = h(start)$.
- After each iteration, the cutoff is increased to the **smallest f -value that exceeded the previous cutoff**.
- The algorithm terminates when the goal is found (optimal) or no nodes remain below ∞ cutoff.

Key Differences

Criterion	DFID	IDA*
Cutoff type	Depth (structural — number of edges from root)	f-cost = $g(n) + h(n)$ (informed uses heuristic knowledge)
Search type	Uninformed (blind)	Informed (uses heuristic h)
Optimality guarantee	Optimal for <i>uniform</i> edge costs only	Optimal for <i>any</i> edge costs, provided h is admissible
Overhead per iteration	Lower — just checks depth $\leq L$	Higher — must compute $f(n) = g(n) + h(n)$ for every node visited
When to use	Unknown/unreliable heuristic, uniform costs	Good admissible heuristic available, non-uniform costs
Space complexity	$O(bd)$ — stores current path	$O(bd)$ — stores current path (same as DFID)
Cutoff progression	$L = 0, 1, 2, \dots$ (linear, fixed step)	$f_{cutoff} = \min\{f(n) \mid n \text{ pruned in last iteration}\}$ (adaptive, data-driven)

Exam summary:

- **DFID** = IDDFS = depth-bounded iterative DFS. Uninformed. Uses depth as cutoff.
- **IDA*** = f -cost-bounded iterative DFS. Informed. Uses $f(n) = g(n) + h(n)$ as cutoff.
- Both algorithms achieve $O(bd)$ space — their key advantage over BFS and A* (which both require $O(b^d)$ space).
- DFID is to BFS what IDA* is to A*: the iterative-deepening, space-efficient version.
- In practice, IDA* is preferred for problems with good heuristics and large state spaces (e.g., the 15-puzzle).

Examination 2020

Note: 2020 has the broadest coverage of search questions — spanning Q2, Q3(b), Q3(c) in Section A and Q7(b) in Section B.

- **Q2 (8.75 marks):** Problem formulation + search space formulation + graph trace (UCS, Greedy, A*)
- **Q3(b) (2 marks):** IDDFS — how it combines BFS + DFS benefits
- **Q3(c) (1.75 marks):** BFS admissibility
- **Q7(b) (2 marks):** A* heuristic — consequences of under/overestimation

Q2(a) — Formulating a Search Problem [1.75 marks]

Question: How can you formulate a search problem? Give examples of data structures that can be used to represent a state-space both in explicit and implicit ways.

Formulating a Search Problem

A search problem is formulated by defining **four components**:

1. **Initial state** — the state in which the agent starts. *Example:* In a route-finding problem, the starting city.
2. **Goal test** — a function $\text{isGoal}(s) \rightarrow \{\text{true}, \text{false}\}$ that determines whether a given state s satisfies the goal condition. Can be an explicit set of goal states or an abstract condition.
3. **Successor function** $\text{Succ}(s) \rightarrow \{(a, s')\}$ — given a state s , returns all (action, resulting-state) pairs reachable by legal actions.
4. **Path cost function** $g(p)$ — assigns a numeric cost to a path p (typically the sum of step costs along the path).

Together, these four components define the **search space** — a directed graph where vertices are states and edges are actions. A **solution** is a path from the initial state to any state satisfying the goal test. An **optimal solution** has the lowest path cost among all solutions.

Representing State-Spaces

1. Explicit Representation (State-Space Graph):

The entire state-space graph is stored in memory before search begins. Each state and its connections are enumerated.

Data structures used:

- **Adjacency list:** For each state, store a list of (neighbor, cost) pairs. Memory: $O(|V| + |E|)$.
- **Adjacency matrix:** A $|V| \times |V|$ matrix where entry $[i][j]$ holds the cost of the edge from state i to state j (or ∞ if no edge). Memory: $O(|V|^2)$.
- **Edge list:** Store all edges as triples (from_state, to_state, cost). Memory: $O(|E|)$.

Limitation: Only practical for **small** state spaces (e.g., the Romania map with 20 cities). For problems like the 8-puzzle ($9!/2 = 181,440$ states) or chess ($\sim 10^{47}$ states), explicit storage is infeasible.

2. Implicit Representation (Successor Function):

The state space is **generated on demand** during search. Only the current state and the successor function are stored; the full graph is never materialised.

Data structures used:

- **State representation:** A compact data structure encoding the current configuration. *Examples:* a tuple (x, y) for grid positions; a 2D array for the 8-puzzle board; a set of boolean variables for satisfiability problems.
- **Successor function (as code):** A function that, given a state, **computes** legal next states algorithmically — no lookup table needed. *Example:* For the Water Jug problem, successors of (x, y) are computed by the 6 pouring/filling rules applied to the current values of x and y .
- **Fringe (frontier):** A queue (FIFO for BFS), stack (LIFO for DFS), or priority queue (for UCS/A*) holding states discovered but not yet expanded. Memory depends on the algorithm, not the state-space size.
- **Closed/explored set:** A hash table or set storing states already expanded, preventing redundant work. Uses $O(|V_{\text{expanded}}|)$ memory.

Advantage: Scales to **enormous** state spaces. The 8-puzzle with 10^5 states is searched without ever storing the full graph. Chess engines use implicit representation — they generate moves on the fly and search to depth 20+ without enumerating all 10^{47} positions.

Q2(b) — Problem Formulations [2 marks]

Question: Give the initial state, goal test, successor function, and cost function for each problem.

(i) Travelling Salesman Problem (TSP)

State: (ℓ, V) where ℓ is the current city and V is the set of unvisited cities.

Initial state: $(\text{start_city}, \{\text{all other } N-1 \text{ cities}\})$.

Goal test: $\ell = \text{start_city}$ **and** $V = \emptyset$.

Successor function: $\text{Travel}(c_i \rightarrow c_j)$ for any $c_j \in V$ (or $c_j = \text{start_city}$ if $V = \emptyset$). Removes c_j from V .

Path cost: Sum of road distances travelled. Goal: find the **minimum-cost** Hamiltonian cycle.

Complexity: $(N-1)!/2$ distinct tours — exponential. This is the canonical NP-hard problem.

(ii) Missionaries & Cannibals

State: (M_L, C_L, B) where $M_L, C_L \in \{0, 1, 2, 3\}$ are missionaries and cannibals on the left bank, and $B \in \{L, R\}$ is the boat position.

Initial state: $(3, 3, L)$.

Goal test: $(0, 0, R)$.

Successor function: Move 1 or 2 people across the river by boat. Valid combinations: $(1M, 0C)$, $(2M, 0C)$, $(0M, 1C)$, $(0M, 2C)$, $(1M, 1C)$. After each move, verify the **safety constraint**: on both banks, $M \geq C$ or $M = 0$.


Path cost: 1 per boat crossing. Goal: **minimum crossings** (known to be 11).

Q2(c) — Graph Traces [5 marks]

Question: Using the search space given (see Table), draw the state-space, then trace UCS, Greedy, and A* from A to G. At each step show which node is being expanded and the content of the fringe.

The 2020 search space is identical to the 2022 and 2023 search space (same 6 nodes A-F plus goal G, same directed edges, same $h(n)$ values). The complete step-by-step traces for all three algorithms on this graph are given in **Section 2.3** (Examination 2023, Part b(ii)). Below is the summary and the state-space diagram.

Edge Table			Heuristic $h(n)$	
State	Next	Cost	Node	$h(n)$
A	B	4	A	8
A	C	1	B	8
B	D	3	C	6
B	E	8	D	5
C	C	0	E	1
C	D	2	F	4
C	F	6	G	0
D	C	2		
D	E	4		
E	G	2		
F	G	8		



Algorithm Results Summary:

Algorithm	Path	Cost	Expansions	Optimal?
UCS	$A \rightarrow C \rightarrow D \rightarrow E \rightarrow G$	9	7	Yes
Greedy	$A \rightarrow C \rightarrow F \rightarrow G$	15	4	No
A*	$A \rightarrow C \rightarrow D \rightarrow E \rightarrow G$	9	5	Yes


For the full step-by-step traces (node expanded, fringe contents at each step, discussion of why Greedy fails and A* succeeds), see **Section 2.3** (Examination 2023). The traces on this graph are among the most important in the syllabus — the same search space appears in 3 exam years (2020, 2022, 2023).

Q3(b) — IDDFS Combines BFS + DFS [2 marks]

Question: Iterative-deepening (ID) is a type of search algorithm said to combine the benefits of breadth-first and depth-first search. Explain, using a diagram or otherwise, how ID achieves these benefits.

The Problem: BFS guarantees the shortest path but needs $O(b^d)$ memory — infeasible for large search spaces. DFS uses only $O(bm)$ memory but may never terminate on infinite branches and finds suboptimal deep solutions.

The IDDFS Solution: Run depth-limited DFS repeatedly with cutoff $L = 0, 1, 2, \dots$ until the goal is found.



L	Nodes visited (DFS order)	Space used
0	A	1 node (just A)
1	A, B, C	2 nodes (A→B, then A→C)
2	A, B, D, E, C, F, I	3 nodes (A→B→D at deepest)
3	A, B, D, G ← found!	3 nodes

How IDDFS inherits from each:

Property	Inherited from BFS	Inherited from DFS
Completeness	✓ BFS explores level by level, never trapped by infinite branches. ID-DFS does the same — depth limits prevent infinite descent.	—
Optimality	✓ BFS finds the shallowest goal. ID-DFS does the same — because it exhausts depth $d-1$ before checking depth d , the first goal found is shallowest.	—
Space	—	✓ DFS uses $O(bm)$ memory — only the current path. IDDFS replicates this: at any moment, only $O(bL)$ nodes are stored (the current path plus siblings). For $b=10, L=10$: ~ 100 nodes vs BFS's 10 billion.
Simple implementation	—	✓ Like DFS, IDDFS uses a simple stack or recursion. The outer loop adds minimal complexity.

Cost of the compromise — Time Overhead: IDDFS re-expands nodes near the root: node A is expanded $d+1$ times, depth-1 nodes are expanded d times, etc. However, the deepest level dominates the node count. The overhead factor is $\frac{b}{b-1}$:

$$\text{Total work} = b^d \left(1 + \frac{1}{b} + \frac{1}{b^2} + \dots \right) \approx b^d \cdot \frac{b}{b-1}$$

For $b = 10$, the overhead is only $\frac{10}{9} \approx 1.11\times$ — an acceptable price for reducing memory from $O(b^d)$ to $O(bd)$.

IDDFS = BFS completeness + BFS optimality + DFS space efficiency. The only trade-off is a modest constant-factor increase in time ($\approx 11\%$ for $b=10$), making it the **preferred uninformed search strategy** for large state spaces with uniform costs.

Q3(c) — BFS Admissibility [1.75 marks]

Question: When is breadth-first search an admissible search strategy? Briefly explain.

What “Admissible” Means

In search, an algorithm is **admissible** if it **always finds the optimal (lowest-cost) solution** when one exists. The term is most commonly applied to heuristics ($h(n) \leq h^*(n)$), but it can also describe complete search algorithms.


When BFS is Admissible

BFS is admissible **when all step costs are equal (uniform)**. Under this condition:

1. BFS explores the state space **level by level**: all nodes at depth 0, then depth 1, then depth 2, ...
2. Since every step costs the same amount (say, 1), the **depth of a node equals its path cost**: $g(n) = \text{depth}(n)$.
3. Therefore, the **first goal node encountered** is at the shallowest depth, which corresponds to the **minimum path cost**.
4. BFS returns this shallowest goal — which is optimal.

When BFS is NOT Admissible

BFS is **not admissible when step costs are non-uniform**. Consider:



BFS would find G_1 first (depth 1, cost 100) and return it as the solution. But the optimal path is $S \rightarrow A \rightarrow G_2$ (depth 2, cost $1 + 1 = 2$). BFS returns the **shallowest** goal, not the **cheapest**. The solution is suboptimal, so BFS is not admissible when edge costs differ.

Formal Statement

BFS is admissible iff all step costs are equal.

When costs are uniform, depth $\propto$ cost, and the shallowest goal is the cheapest. When costs vary, BFS returns the shallowest goal regardless of cost — this may be suboptimal. For non-uniform costs, use **UCS** (Uniform-Cost Search) instead — it expands by $g(n)$ rather than depth, guaranteeing optimality for any non-negative edge costs.

Q7(b) — A* Heuristic: Underestimate vs Overestimate [2 marks]

Question: What will happen if h' underestimates h in the A* algorithm? What will happen if h' overestimates h ?

Notation clarification: In standard A* notation, $h(n)$ is the heuristic estimate of the cost from n to the goal, and $h^*(n)$ is the *true* optimal cost. The question uses h' for the heuristic and h for the true cost. We adopt the standard convention below.

Case 1: Heuristic UNDERESTIMATES ($h(n) \leq h^*(n)$ — Admissible)

When the heuristic never overestimates the true cost to the goal (it is **admissible**):

Consequence	Explanation
A* is optimal — it will find the lowest-cost solution.	The proof: A* expands nodes in order of $f(n) = g(n) + h(n)$. Since $h(n) \leq h^*(n)$, we have $f(n) \leq g(n) + h^*(n) = \text{true total cost through } n$. The first time A* expands a goal node G , $f(G) = g(G) + 0 = g(G)$. For any unexpanded node n on a cheaper path to a different goal, $f(n) \leq \text{true cost of that cheaper path} < g(G)$. So A* would have expanded n before G — contradiction. Therefore the first expanded goal must be optimal.
May expand more nodes than with a perfect heuristic.	A severe underestimate (e.g., $h(n) = 0$ for all n , which reduces A* to UCS) leads to more node expansions. The closer h is to h^* , the fewer nodes A* expands.
Still complete — will find a solution if one exists.	Admissibility does not affect completeness. A* with any heuristic eventually explores the entire (finite) search space if needed.

Example — admissible (underestimate): $h(\text{Arad}) = 366 \leq h^*(\text{Arad}) = 418$. A* finds the optimal path $\text{Arad} \rightarrow \text{Sibiu} \rightarrow \text{Rimnicu} \rightarrow \text{Pitesti} \rightarrow \text{Bucharest}$ (cost 418).

Case 2: Heuristic OVERESTIMATES ($h(n) > h^*(n)$ for some n)

When the heuristic **overestimates** the true cost for at least one node (it is **not admissible**):

Consequence	Explanation
A* may NOT be optimal — it can return a suboptimal solution.	If $h(n)$ overestimates the remaining cost through node n , then $f(n)$ is inflated. A* may postpone expanding n in favor of nodes whose f -values are lower but lie on suboptimal paths. A suboptimal goal may be expanded and returned before the optimal path is discovered.
The optimal path may be “hidden.”	A* considers nodes in order of f -value. An overestimating heuristic makes good nodes look worse than they are (artificially high f), so they sit deeper in the fringe while A* explores suboptimal-looking nodes with deceptively low f .
Still complete — overestimation does not affect completeness.	A* will still find <i>a</i> solution (the search is still systematic). It just may not be the <i>best</i> one.

Example — inadmissible (overestimate): Suppose $h(\text{Fagaras}) = 200$, but the true cost from Fagaras to Bucharest is $h^*(\text{Fagaras}) = 99$. Then $f(\text{Fagaras}) = 239 + 200 = 439$, while the path through Pitesti has $f(\text{Pitesti}) = 317 + 101 = 418$. The Fagaras path looks worse and is explored later. If $\text{Sibiu} \rightarrow \text{Fagaras} \rightarrow \text{Bucharest}$ (cost $239 + 99 = 338$) were actually optimal but h made it look worse, A* might return the Pitesti route (cost 418) — suboptimal.

Summary

Property	h underestimates (admissible)	h overestimates (inadmissible)
Optimality	Guaranteed ✓	Not guaranteed ×
Node expansions	More than with exact h	Unpredictable — may expand fewer nodes
Completeness	Guaranteed ✓	Guaranteed ✓
Example	$h(n) = 0$ (UCS — slow but optimal)	$h(n) > h^*(n)$ (fast but potentially wrong)
Design rule	Always make h admissible	Never intentionally overestimate

Exam key point: The **admissibility condition** $h(n) \leq h^*(n)$ for all n is the **single most important property** of a heuristic in A*. It is the necessary and sufficient condition for A* optimality (on tree-search; on graph-search, consistency/monotonicity is also required for optimality without reopening nodes). An admissible heuristic guarantees you won't miss the best path; a well-chosen admissible heuristic (close to h^*) makes the search efficient.